

STDCL

A Simplified C Interface for OpenCL

revision 1.2

Copyright © 2009-2011 Brown Deer Technology, LLC

Verbatim copying and distribution of this entire document is permitted in any medium, provided this notice is preserved.

Contents

- [Name](#)
 - [Version](#)
 - [Synopsis](#)
 - [Description](#)
 - [Application Programming Interface \(API\)](#)
 - [Default Contexts](#)
 - [Dynamic CL Program Loader](#)
 - [Memory Management](#)
 - [Kernel Management](#)
 - [Synchronization](#)
 - [Environment Variables](#)
 - [Examples](#)
 - [Manual Pages](#)
-

Name

STDCL - Standard Compute Layer Interface

Version

STDCL_VERSION_STR

STDCL_VERSION_HEX

Synopsis

`#include <stdcl.h>`

Link with `-lstdcl`.

Default Contexts

`stddev, stdcpu, stdgpu, stdrpu,`
`clgetndev()`

Dynamic CL Program loader

`clopen(), clsym(), clclose(), clbuild()`

Memory Management

`clmalloc(), clmrealloc(), clfree(), clglmalloc(),`
`clmctl(), clmctl_va(),`
`clmattach(), clmdetach(), clsizeofmem(), clmsync(), clglmsync()`

Kernel Management

`clndrange_init1d(), clndrange_init2d(), clndrange_init3d(),`
`clarg_set(), clarg_set_local(), clarg_set_global(),`
`clfork()`

Synchronization

`clflush(), clwait()`

Environment Variables

`STDDEV, STDCPU, STDGPU, STDRPU,`
`STD[DEV|CPU|GPU|RPU]_PLATFORM_NAME,`
`STD[DEV|CPU|GPU|RPU]_MAX_NDEV,`
`STD[DEV|CPU|GPU|RPU]_LOCK`

Description

OpenCL provides a host-side API that allows the careful management of memory and processes on heterogeneous computing platforms. The level of control is more typically reserved for conventional operating systems (memory management, process management, synchronization, etc.). Although this granularity of control is necessary to support the expansive industry objectives for which OpenCL was designed, the granularity of control and verbose nature of the API proves to be tedious within the context of typical software application development. The steps required for a simple Hello World OpenCL program are tedious and repetitive from a programmer's perspective. Moreover, some semantics introduced by OpenCL have more natural and familiar constructs within traditional UNIX programming that can greatly simplify the use of the API and prove more efficient. As an example, opaque memory buffers are more naturally managed as memory allocations; modern UNIX-like operating systems are more than capable of employing

memory virtualization sufficient to allow control over memory consistency.

STDCL provides a simplified C interface to OpenCL designed in a style familiar to traditional UNIX/C programmers. The design and implementation of STDCL is inspired by familiar APIs designed for different purposes, e.g., `stdio.h` (for default contexts), `dlopen` (for managing OpenCL kernels), `malloc` (as a replacement for creating opaque memory buffers), and `fork` (as a replacement to "enqueueing commands on the command queue"). In every detail, the approach is to avoid introducing new inventive syntax and semantics in favor of exploiting permutations of more familiar syntax and semantics from traditional UNIX. Whether the effort succeeds is for the programmer to decide.

Application Programming Interface (API)

The STDCL interface provides support for [default contexts](#), [dynamic CL program loader](#), [memory management](#), [kernel execution](#), and [asynchronous operations](#). In addition, [environment variables](#) provide run-time control over certain aspects of the interface. The STDCL interface is discussed in detail below.

Default Contexts

STDCL provides several default contexts similar to the default I/O streams provided by `stdio`. These default contexts are defined to include the most typical use-cases. Each default context is of type `CLCONTEXT`, which is defined as a superset of the OpenCL type `cl_context`. The following default contexts are provided:

`CLCONTEXT* stddev;`

All devices for a given platform supported by the OpenCL API.

`CLCONTEXT* stdcpu;`

All multi-core CPU processors for a given platform supported by the OpenCL API.

`CLCONTEXT* stdgpu;`

All many-core GPU processors for a given platform supported by the OpenCL API.

`CLCONTEXT* stdrpu;`

All reconfigurable processors for a given platform supported by the OpenCL API.

`cl_uint clgetndev(CLCONTEXT* cp);`

This call returns the number of devices in the CL context `cp`.

Dynamic CL Program Loader

STDCL provides a convenient interface for dynamically loading CL programs and accessing OpenCL kernels. The functions `clopen()`, `clsym()` and `clclose()` are designed to mirror the semantics of the more familiar functions `dlopen()`, `dlsym()` and `dlclose()` used to access the Linux dynamic loader. The following functions are provided for dynamically loading CL programs and accessing OpenCL kernels:

`void* clopen(CLCONTEXT* cp, const char* filename, int flags);`

This call opens a file containing the source or binary program defining one or more OpenCL kernels and

performs the steps necessary to create and build the OpenCL program object. A handle is returned that can be used in subsequent calls to access the actual kernels in the program. The handle is valid within the CLCONTEXT specified by **cp**.

If **filename** is a NULL pointer then a handle to the OpenCL program(s) embedded in the host program executable is returned. (See the tool `clld` for a description of how to embed OpenCL source and binary programs into a host program executable.)

The **flags** argument allows control over the behavior of the function. The flag `CLLD_NOW` instructs the call to perform all of the steps involved with creating and building the program; the flag `CLLD_LAZY` instructs the call to defer these steps until the handle is first used. The call accepts a flag set to 0 in which case the default behavior (`CLLD_NOW`) is used.

If the **flags** argument `CLLD_NOBUILD` is used the compilation and build process is deferred, and a subsequent call to `clbuild()` must be used for the returned handle to reflect a valid (compiled and built) program. This flag is useful when the user needs to pass in compiler options, which can be done with the `clbuild()` call.

By default the following compiler options will always be passed to the low-level OpenCL calls:

```
-D __STDCL__  
-D __CPU__ | __GPU__  
-D __AMD__ | __NVIDIA__ | __coprthr__  
-I $(CWD)
```

```
void* clsopen( CLCONTEXT* cp, const char* srcstr, int flags );
```

This call behaves exactly like `clopen()` with the exception that instead of providing the name of a file containing the OpenCL kernel source, the kernel code may be provided directly as a string. This call can be useful within schemes where custom kernel source is generated at run-time.

```
cl_kernel clsym( CLCONTEXT* cp, void* handle, const char* symbol, int  
flags );
```

This call takes a **handle** returned from a call to `clopen()` and returns the OpenCL kernel specified by **symbol**. The OpenCL kernel is created within the CLCONTEXT specified by **cp**.

The argument **flags** allows control over the behavior of the function. The flag `CLLD_NOW` instructs the call to perform all of the steps involved with creating the kernel; the flag `CLLD_LAZY` instructs the call to defer these steps until the kernel is first used. The call accepts a flag set to 0 in which case the default behavior (`CLLD_NOW`) is used.

```
int clclose( CLCONTEXT* cp, void* handle );
```

This call decrements the reference count on the associated handle. If the reference count drops to zero then the associated OpenCL program source or binary is unloaded and the associated file is closed. Under normal usage this call is used to safely release the OpenCL programs created by a call to `clopen()`.

```
void* clbuild( CLCONTEXT* cp, void* handle, char* options, int flags );
```

This call is used following a call to `clopen()` or `clsopen()` with the `CLLD_NOBUILD` flag. Calling `clbuild()` will complete the process of compiling and building the kernel program. This call accepts user-specified compiler options. The handle passed in must be a valid handle created by a call to

`clopen()` or `clsopen()` with the `CLLD_NOBUILD` flag. Calling `clbuild()` will complete the porocess of compilinnng and building the kernel program. This call accepts user-specified compiler options. The handle passed in must be a valid handle created by a call to `clopen()` or `clsopen()` with the `CLLD_NOBUILD` flag. In addition to the user defined compiler options, the standard compiler options described for `clopen()` are also passed to low-level OpenCL calls.

Memory Management

STDCL provides functions for allocating and managing memory that may be shared between the host and OpenCL co-processor devices. Memory may be allocated with `clmalloc()` and used transparently as the global memory for kernel execution on a OpenCL device. The programmer uses a single pointer representing the allocated memory which may be re-attached to various CL contexts using `clmattach()` and `clmdetach()`. Memory consistency can be maintained using the `clmsync()` function which synchronizes memory between the host and OpenCL co-processor devices. The following functions are provided for OpenCL memory management.

```
void* clmalloc( CLCONTEXT* cp, size_t size, int flags);
```

This call allocates memory suitable for sharing between OpenCL co-processor devices within a CL context. The size of the allocation is specified in bytes. The memory is not cleared. The last argument is used to pass flags to control the behavior of function. The flag `CL_MEM_DETACHED` may be used to allocate memory that is not attached to a CL context in which case `cp` must be 0. If `flags` is 0 the default behavior is to allocate memory attached to a specified CL context.

```
void* clmrealloc( CLCONTEXT* cp, void* ptr, size_t size, int flag);
```

This call re-allocates memory suitable for sharing between OpenCL co-processor devices within a CL context and may be used to change the size of an existing allocation. The `ptr` argument must be a valid memory allocation returned by either `clmalloc()`, or a previous call to `clmrealloc()`. The size of the allocation is specified in bytes. The memory is not cleared. The last argument is used to pass flags to control the behavior of function. The flag `CL_MEM_DETACHED` may be used to allocate memory that is not attached to a CL context in which case `cp` must be 0. If `flags` is 0 the default behavior is to allocate memory attached to a specified CL context.

```
void* clglmalloc( CLCONTEXT* cp, cl_GLuint glbufobj, int flag);
```

This call allocates memory suitable for sharing between OpenCL co-processor devices within a CL context based on an existing OpenGL memory buffer `glbufobj`. The size of the allocation is implied by the OpenGL buffer size. The memory is not cleared. The last argument is used to pass flags to control the behavior of function. The flag `CL_MEM_DETACHED` may be used to allocate memory that is not attached to a CL context in which case `cp` must be 0. If `flags` is 0 the default behavior is to allocate memory attached to a specified CL context.

```
void clfree( void* ptr);
```

This call frees memory allocated with `clmalloc()`. The memory specified by `ptr` can be either attached or detached from a CL context. Calling `clfree()` with `ptr` equal to 0 is considered an error.

```
size_t clsizeofmem(void* ptr);
```

This call returns the size in bytes of the memory allocated with `clmalloc()`.

```
int clmctl( void* ptr, int op, ... );
```

```
int clmctl_va( void* ptr, int op, va_list );
```

These calls provide generalized control over a device-sharable memory allocation and differ only in the way optional arguments are passed in. the **ptr** argument is a pointer to device-sharable memory returned by any of the calls `clmmalloc()`, `clmrealloc()`, or `clglmalloc()`. The following operations for the `op` argument are presently valid:

CL_MCTL_SET_IMAGE2D

Mark the allocation to be of OpenCL image2d_t type. Optional arguments are:

size_t sz0

Image width

size_t sz1

Image height

size_t sz2

Should be set to 0

cl_image_format* fmt

(Optional) pointer to image format struct, ignored if set to NULL

```
cl_event clmsync( CLCONTEXT* cp, unsigned int devnum, void* ptr, int flags );
```

This call is used to synchronize memory between the host platform and OpenCL co-processor devices. The memory specified by **ptr** must have been allocated by `clmmalloc()` and associated with a CL context.

The behavior of `clmsync()` is controlled by the **flags** argument which must be set with either `CL_MEM_HOST` or `CL_MEM_DEVICE`. These flags are mutually exclusive and it is an error to set both or none. In addition the flags `CL_EVENT_WAIT` and `CL_EVENT_NOWAIT` control the blocking behavior for the call. For a blocking call the flag `CL_EVENT_NORELEASE` may be specified to prevent the call from releasing OpenCL events created as a result of the call. If the flag `CL_EVENT_NORELEASE` is specified, the programmer is responsible for releasing the returned event with the OpenCL call `clReleaseEvent()`.

The following examples demonstrate typical uses of `clmsync()`:

Non-blocking sync to device memory:

```
clmsync(stdgpu,0,ptr,CL_MEM_DEVICE|CL_EVENT_NOWAIT);
```

Non-blocking sync to host memory:

```
clmsync(stdgpu,0,ptr,CL_MEM_HOST|CL_EVENT_NOWAIT);
```

Blocking sync to device memory:

```
clmsync(stdgpu,0,ptr,CL_MEM_DEVICE|CL_EVENT_WAIT);
```

Blocking sync to host with release of event:

```
clmsync(stdgpu,0,ptr,CL_MEM_HOST|CL_EVENT_WAIT);
```

```
cl_event clglmsync( CLCONTEXT* cp, unsigned int devnum, void* ptr, int flags );
```

This call is used to sync memory between device-sharable memory and OpenGL buffers.

The flags argument must be set to either `CL_MEM_CLBUF` or `CL_MEM_GLBUF` to define the

destination of the sync operation.

```
int clmattach( CLCONTEXT* cp, void* ptr );
```

This call is used to attach memory allocated by `clmalloc()` to a CL context. In order to change the attachment of memory from one CL context to another, the memory must first be unattached using a call to `clmdetach()`. It is an error to call with a **ptr** to memory that is already attached to a CL context.

```
int clmdetach( void* ptr );
```

This call is used to detach memory from a CL context. The memory must have been allocated by `clmalloc()`.

Kernel Management

STDCL provides simplified interfaces for setting up the index-space and arguments for kernel execution. Executing a kernel on an OpenCL co-processor device is supported using `clfork()` which allows blocking and non-blocking execution behavior. The following functions are provided for OpenCL kernel management.

```
clndrange_t clndrange_init1d( gtoff0,gtsz0,ltsz0 );  
clndrange_t clndrange_init2d( gtoff0,gtsz0,ltsz0, gtoff1,gtsz1,ltsz1 );  
clndrange_t clndrange_init3d( gtoff0,gtsz0,ltsz0, gtoff1,gtsz1,ltsz1,  
gtoff2,gtsz2,ltsz2 );
```

The `clndrange_init*`() functions are used to *initialize* a variable of type `clndrange_t` used to store the OpenCL index-space over which a kernel is to execute. These functions will be implemented as macros to allow for struct initialization in C. The arguments **gtoff**, **gtsz** and **ltsz** represent the global offset, global size and local size of the index-space for a given dimension, respectively. As an example, the following initializes a two dimensional OpenCL NDRange with no offsets over a global index space of size 512 by 2048 with a local work group size of 4 by 16:

```
clndrange_t ndr = clndrange_init2d( 0,512,4 0,2048,16 );
```

```
void clarg_set( CLCONTEXT* cp, cl_kernel krn, unsigned int argnum, Tn arg  
);
```

This call is used to set the argument of an OpenCL kernel for arguments of intrinsic non-pointer type that are to be passed by value. The size of the argument is inferred from the type of the argument and may be a vector type, e.g., `cl_float4`.

```
void clarg_set_global( CLCONTEXT* cp, cl_kernel krn, unsigned int argnum,  
void* ptr );
```

This call is used to set the argument of an OpenCL kernel for arguments that are pointers to global memory as defined in the OpenCL specification. The memory must have been allocated by `clmalloc()` in the appropriate CL context of the kernel.

```
void clarg_set_local( CLCONTEXT* cp, cl_kernel krn, unsigned int argnum,  
size_t sizeb );
```

This call is used to set the argument of an OpenCL kernel for arguments that are pointers to local memory as defined in the OpenCL specification. Local memory of size **sizeb** bytes will be allocated for use by the OpenCL kernel.

```
cl_event clfork( CLCONTEXT* cp, unsigned int devnum, cl_kernel krn,  
clndrange* ndr, int flags );
```

This call is used to execute a kernel on the OpenCL co-processor device specified by **devnum**. The arguments for the kernel must be set prior to the call to `clfork()` using the `clarg_set*()` functions described above. The kernel is executed over an index-space of work-items defined by **ndr**.

The behavior of `clfork()` may be controlled using the flags `CL_EVENT_WAIT` or `CL_EVENT_NOWAIT`. Specifying the flag `CL_EVENT_NOWAIT` will cause `clfork()` to return immediately. Specifying the flag `CL_EVENT_WAIT` will cause `clfork()` to block until the kernel execution is complete. Including the flag `CL_EVENT_NORELEASE` will prevent the event associated with the kernel execution to be released for blocking calls to `clfork()`. If the flag `CL_EVENT_NORELEASE` is specified the programmer is responsible for releasing the returned event with the OpenCL call `clReleaseEvent()`.

The following examples demonstrate typical uses of `clfork()`:

Blocking execution of a kernel on device number 0 automatically releasing the associated event:

```
clfork( stdgpu, 0, my_krn, &ndr, CL_EVENT_WAIT );
```

Non-blocking execution of a kernel on device number 2 automatically releasing the associated event:

```
clfork( stdgpu, 2, my_krn, &ndr, CL_EVENT_NOWAIT );
```

Synchronization

STDCL provides functions for synchronization to manage the inherently asynchronous operations enabled by OpenCL per device within each CL context.

```
int clflush( CLCONTEXT* cp, unsigned int devnum, int flags );
```

This call is used to flush all commands enqueued in the command queue associated with the OpenCL device specified by the device number **devnum** within the specified CL context. For typical OpenCL implementations this is necessary to force the execution of commands without blocking on the host. A call to `clflush()` is non-blocking and will return immediately. At present the argument **flags** should be set to 0.

```
cl_event clwait( CLCONTEXT* cp, unsigned int devnum, int flags );
```

This call is used to block on the completion of all commands enqueued in the command queue associated with the OpenCL device specified by the device number **devnum** within the specified CL context.

The **flags** argument is used to control the behavior of the call as follows. The flag `CL_KERNEL_EVENT` will cause the call to block on completion of all enqueued kernel events enqueued by calls to `clfork()`. the flag `CL_MEM_EVENT` will cause the call to block on completion of all enqueued memory events enqueued by call to `clmsync()`. The flags `CL_KERNEL_EVENT` and `CL_MEM_EVENT` may be combined in a single call. Including the flag `CL_EVENT_NORELEASE` will prevent all OpenCL events to be released before `clwait()` returns. If the flag `CL_EVENT_NORELEASE` is specified the programmer is responsible for releasing all events with the OpenCL call `clReleaseEvent()`.

The following examples demonstrate typical uses of `clwait()`:

Block on completion of all kernel execution events on OpenCL device number 0 releasing all events:

```
clwait( stdgpu, 0, CL_KERNEL_EVENT );
```


Block on completion of all memory events on OpenCL device number 2 releasing all events:

```
clwait( stdgpu, 2, CL_MEM_EVENT );
```

Block on completion of all kernel and memory events on OpenCL device number 2 releasing all events:

```
clwait( stdgpu, 2, CL_ALL_EVENT );
```

Environment Variables

The run-time behavior of STDCL can be controlled using environment variables as follows.

STDDEV, STDCPU, STDGPU, STDRPU

Each default CL context is can be controlled by the associated environment variable. A value of 0 or 1 will disable or enable the context, respectively. The default behavior is to attempt to enable the context if valid devices are available within the selected platform.

STD[DEV|CPU|GPU|RPU]_PLATFORM_NAME

Set the platform name for the desired platform to be used for a given context. If none is provided or the specified platform is unavailable, a default will be selected.

STD[DEV|CPU|GPU|RPU]_MAX_NDEV

Set the maximum number of devices for a given context regardless of whether more devices exist for the selected platform. If there is an insufficient number of devices, the maximum available will be provided.

STD[DEV|CPU|GPU|RPU]_LOCK

Set a lock ID for the process to enforce exclusive access to the provided devices across all processes run with the same lock ID. This feature is primarily useful to ensure the multiple MPI processes on a multi-GPU platform are each given exclusive access to a GPU with no requirement on the application itself.

Examples

The following example shows the use of STDCL for a simple program that adds two vectors on a GPU or a CPU:

```
/* example #1 */

#include <stdio.h>
#include <strings.h>
#include <stdcl.h>

#define SIZE 1024

int main()
{
    int i;

    CLCONTEXT* cp = (stdgpu)? stdgpu : stdcpu;
```

```

void* clh = clopen(cp, "add_vec.cl",CLLD_NOW);
cl_kernel k_addvec = clsym(cp, clh, "addvec_kern", CLLD_NOW);

float* aa = (float*)clmalloc(cp,SIZE*sizeof(float),0);
float* bb = (float*)clmalloc(cp,SIZE*sizeof(float),0);
float* cc = (float*)clmalloc(cp,SIZE*sizeof(float),0);

for(i=0;i<SIZE;i++) {
    aa[i] = 111.0f * i;
    bb[i] = 222.0f * i;
}

bzero(cc,SIZE*sizeof(float));

clndrange_t ndr = clndrange_init1d(0,SIZE,64);

clmsync(cp,0,aa,CL_MEM_DEVICE|CL_EVENT_NOWAIT);
clmsync(cp,0,bb,CL_MEM_DEVICE|CL_EVENT_NOWAIT);

clarg_set_global(cp,k_addvec,0,aa);
clarg_set_global(cp,k_addvec,1,bb);
clarg_set_global(cp,k_addvec,2,cc);

clfork(cp,0,k_addvec,&ndr,CL_EVENT_NOWAIT);

clmsync(cp,0,cc,CL_MEM_HOST|CL_EVENT_NOWAIT);

clwait(cp,0,CL_MEM_EVENT|CL_KERNEL_EVENT);

for(i=0;i<SIZE;i++) printf("%f %f %f\n",aa[i],bb[i],cc[i]);

if (aa) clfree(aa);
if (bb) clfree(bb);
if (cc) clfree(cc);

clclose(cp,clh);
}

```

The following example shows the use of STDCL for a simple program that adds two vectors on two GPU:

```

/* example #2 */

#include <stdio.h>
#include <strings.h>
#include "stdcl.h"

#define SIZE 1024

int main()
{
    int i,n;

    CLCONTEXT* cp = stdgpu;

```

```

void* clh = clopen(cp, "add_vec.cl",CLLD_NOW);
cl_kernel k_addvec = clsym(cp, clh, "addvec_kern", CLLD_NOW);

float* aa[2];
float* bb[2];
float* cc[2];

aa[0] = (float*)clmalloc(cp,SIZE*sizeof(float)/2,0);
aa[1] = (float*)clmalloc(cp,SIZE*sizeof(float)/2,0);
bb[0] = (float*)clmalloc(cp,SIZE*sizeof(float)/2,0);
bb[1] = (float*)clmalloc(cp,SIZE*sizeof(float)/2,0);
cc[0] = (float*)clmalloc(cp,SIZE*sizeof(float)/2,0);
cc[1] = (float*)clmalloc(cp,SIZE*sizeof(float)/2,0);

for(i=0;i<SIZE/2;i++) {
    aa[0][i] = 111.0f * i;
    aa[1][i] = 111.0f * (SIZE/2 + i);
    bb[0][i] = 222.0f * i;
    bb[1][i] = 222.0f * (SIZE/2 + i);
}

bzero(cc[0],SIZE*sizeof(float));
bzero(cc[1],SIZE*sizeof(float));

clndrange_t ndr = clndrange_init1d(0,SIZE/2,64);

clmsync(cp,0,aa[0],CL_MEM_DEVICE|CL_EVENT_NOWAIT);
clmsync(cp,1,aa[1],CL_MEM_DEVICE|CL_EVENT_NOWAIT);
clmsync(cp,0,bb[0],CL_MEM_DEVICE|CL_EVENT_NOWAIT);
clmsync(cp,1,bb[1],CL_MEM_DEVICE|CL_EVENT_NOWAIT);

clarg_set_global(cp,k_addvec,0,aa[0]);
clarg_set_global(cp,k_addvec,1,bb[0]);
clarg_set_global(cp,k_addvec,2,cc[0]);

clfork(cp,0,k_addvec,&ndr,CL_EVENT_NOWAIT);

clmsync(cp,0,cc[0],CL_MEM_HOST|CL_EVENT_NOWAIT);

clflush(cp,0,0);

clarg_set_global(cp,k_addvec,0,aa[1]);
clarg_set_global(cp,k_addvec,1,bb[1]);
clarg_set_global(cp,k_addvec,2,cc[1]);

clfork(cp,1,k_addvec,&ndr,CL_EVENT_NOWAIT);

clmsync(cp,1,cc[1],CL_MEM_HOST|CL_EVENT_NOWAIT);

clflush(cp,1,0);

clwait(cp,0,CL_MEM_EVENT|CL_KERNEL_EVENT);
clwait(cp,1,CL_MEM_EVENT|CL_KERNEL_EVENT);

```

```

for(i=0;i<SIZE/2;i++) printf("%f %f %f\n",aa[0][i],bb[0][i],cc[0][i]);
for(i=0;i<SIZE/2;i++) printf("%f %f %f\n",aa[1][i],bb[1][i],cc[1][i]);

if (aa[0]) clfree(aa[0]);
if (aa[1]) clfree(aa[1]);
if (bb[0]) clfree(bb[0]);
if (bb[1]) clfree(bb[1]);
if (cc[0]) clfree(cc[0]);
if (cc[1]) clfree(cc[1]);

clclose(cp,clh);
}

```

Manual Pages

overview	stdcl(3)
dynamic loader	clopen(3) , clsym(3) , clclose(3)
memory management	clmalloc(3) , clfree(3) , clsizeofmem(3) , clmsync(3) , clmattach(3) , clmdetach(3)
kernel management	clndrange_init1d(3) , clndrange_init2d(3) , clndrange_init3d(3) , clarg_set(3) , clarg_set_global(3) , clarg_set_local(3) , clfork(3)
asynchronous operations	clflush(3) , clwait(3)

STDLC (3)

Standard Compute Layer (CL) Manual

STDLC (3)

NAME

stdcl - standard compute layer (CL) library functions

SYNOPSIS

```
#include <stdcl.h>
```

```
CLCONTEXT* stddev;
CLCONTEXT* stdcpu;
CLCONTEXT* stdgpu;
CLCONTEXT* stdrpu;
```

Link with -lstdcl.

DESCRIPTION

The standard compute layer (CL) library (libstdcl) provides a simplified interface to OpenCL designed to support the most typical use-cases in a style inspired by familiar and traditional UNIX APIs for C programming.

libstdcl provides managed OpenCL contexts identified with a context

pointer that is generally provided as an argument to library functions that transparently manage OpenCL constructs such as contexts, devices, memory, kernels and events in a manner that simplifies their use.

Default Contexts

libstdcl provides several default contexts similar to the default I/O streams provided by stdio. the following default contexts are provided:

stddev All devices for a given platform supported by the OpenCL API.

stdcpu All multi-core CPU processors for a given platform supported by the OpenCL API.

stdgpu All many-core GPU processors for a given platform supported by the OpenCL API.

stdrpu All reconfigurable processors for a given platform supported by the OpenCL API.

Dynamic CL Program Loader

libstdcl provides a convenient interface for dynamically loading CL programs and accessing CL kernels. Using the tool clld CL program source and binary files can be embedded within special ELF sections linked against other object files on the host platform to generate a single executable. The set of functions `clopen()`, `clsym()`, `clclose()` provide a convenient interface capable of dynamically loading CL programs embedded within the executable as well as from an external file. CL programs.

Memory Management

libstdcl provides functions for allocating and managing memory that may be shared between the host and CL co-processor devices. Memory may be allocated with `clmalloc()` and used transparently as the global memory for kernel execution on a CL device. The programmer uses a single pointer representing the allocated memory which may be re-attached to various CL contexts using `clmattach()` and `clmdetach()`. Memory consistency can be maintained using the `clmsync()` function which synchronizes memory between host and CL co-processor device.

Kernel Execution

libstdcl provides simplified interfaces for setting up the index-space and arguments for kernel execution. Executing a kernel on a particular CL co-processor device is supported using `clfork()` which allows blocking and non-blocking execution behavior.

Management of Asynchronous Operations

libstdcl provides event management per device within each context to simplify the management of asynchronous multi-device operations. The function `clwait()` can be used to block on selected events within one of

several per device event lists managed transparently.

EXAMPLE

The following example shows a very simple program for calculating the outer product of two vectors using a GPU:

```
#include <stdcl.h>

int main() {

    int n = 1024;

    cl_float* aa = (cl_float*)clmalloc(stdgpu,n,0);
    cl_float* bb = (cl_float*)clmalloc(stdgpu,n,0);
    cl_float* cc = (cl_float*)clmalloc(stdgpu,n,0);

    /* initialize aa and bb */

    void* h = clopen(stdgpu,"outer_prod_kern.cl",0);
    cl_kernel krn = clsym(stdgpu,h,"outer_prod_kern");

    clndrange_t ndr = clndrange_init1d(0,n,4);

    clarg_set(stdgpu,krn,0,n);
    clarg_set_global(stdgpu,krn,1,aa);
    clarg_set_global(stdgpu,krn,2,bb);
    clarg_set_global(stdgpu,krn,3,cc);

    clfork(stdgpu,0,krn,ndr,CL_EVENT_NOWAIT);

    clmsync(stdgpu,0,cc,CL_EVENT_NOWAIT);

    clwait(stdgpu,0,CL_ALL_EVENTS|CL_EVENT_RELEASE);

    clclose(h);

    clfree(aa);
    clfree(bb);
    clfree(cc);

}
```

ENVIRONMENT

Executables that use the libstdcl library are affected by environment variables that control the behavior of the API. The environment variables STDDEV, STDCPU, STDGPU, STDRPU may be set to 0 or 1 to disable or enable the context, respectively. Additional environment variables are used to control the behavior of the context.

The environment variables STDDEV_PLATFORM_NAME, STDCPU_PLATFORM_NAME, STDGPU_PLATFORM_NAME, STDRPU_PLATFORM_NAME can be used to set the name of the preferred OpenCL platform to use for the context.

As an example, the following would force the stdgpu context to use the AMD APP platform:

```
setenv STDGPU_PLATFORM_NAME AMD
```

The environment variable `STDDEV_MAX_NDEV`, `STDCPU_MAX_NDEV`, `STDGPU_MAX_NDEV`, `STDRPU_MAX_NDEV` can be used to set the maximum number of devices provided by the respective context regardless of the number of devices available.

The environment variables `STDDEV_LOCK`, `STDCPU_LOCK`, `STDGPU_LOCK`, `STDRPU_LOCK` can be used to setup exclusive locking of devices across processes with matching lock values. In combination with setting the maximum number of devices in the context, the lock can be used to achieve efficient device management between related processes.

As an example, the following would allow the individual MPI processes on a given node to have exclusive access to a single GPU provided that the number of MPI processes per node did not exceed the number of available GPUs:

```
setenv STDGPU_MAX_NDEV 1; setenv STDGPU_LOCK $(MPI_JOB_ID)
```

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3. There is NO WARRANTY to the extent permitted by law.

SEE ALSO

`clld(1)`, `clopen(3)`, `clsym(3)`, `clclose(3)`, `clmalloc(3)`, `clmsync()`, `clfork(3)`, `clwait(3)`

"libstdcl-1.2"

2011-6-13

STDLC(3)

CLOPEN(3)

Standard Compute Layer (CL) Manual

CLOPEN(3)

NAME

`clopen`, `clsym`, `clclose`, `clerror`, `claddr` - programming interface to dynamic CL loader

SYNOPSIS

```
#include <stdcl.h>
```

```
void* clopen( CLCONTEXT* cp, const char* filename, int flags);
```

```

void* clbuild( CLCONTEXT* cp, void* handle, char* uopts, int flags);

cl_kernel clsym( CLCONTEXT* cp, void* handle, const char* symbol, int
flags);

int claddr( CLCONTEXT* cp, void* addr, CL_info* info);

char* clerror( void );

int clclose( CLCONTEXT* cp, void* handle);

Link with -lstdcl.

```

DESCRIPTION

The functions `clopen()`, `clsym()`, `clclose()`, and `clerror()` implement an interface for dynamically loading compute layer (CL) kernels.

The function `clopen()` loads the CL source or binary program file named by the NULL-terminated string `filename` and returns an opaque handle that may be used as a reference in subsequent calls. If `filename` is a NULL pointer then a handle for the main program executable is returned. If the flag `CLLD_NOBUILD` is used then the CL source will not be compiled. This may be used to defer the compilation and build to the `clbuild()` call which can accept compiler options.

The function `clbuild()` is used for deferred compilation and allows compiler options to be passed. The combination of `clopen()` with the flag `CLLD_NOBUILD` and `clbuild()` is equivalent to a standard `clopen()` call.

The function `clsym()` takes a handle to a CL source or binary program and a NULL-terminated symbol name and returns the associated CL kernel. A CL context pointer must be specified to identify the appropriate CL kernel to return. If `handle` is NULL then all CL programs loaded into the specified CL context are searched.

The function `clclose()` decrements the reference count on the associated handle. If the reference count drops to zero then the CL program is unloaded. The function `clclose()` returns the reference count on success and -1 on error.

The function `clerror()` returns a human readable string describing the most recent error that has occurred as a result of a call to any of the functions `clopen()`, `clsym()`, `clclose()` since the last call to `clerror()`. If no error has occurred NULL is returned.

The function `claddr()` takes as an argument a CL kernel and tries to resolve the name and file where it is located. Information is returned in the `cl_kernel_info` structure:

```

struct cl_kernel_info {
    const char* cli_fname;
    CLCONTEXT* cli_cp;
    unsigned int cli_devnum;
    const char* cli_kname;
};

```



```
};
```

If no matching kernel is found the fields are set to NULL. claddr() returns zero on error and non-zero on success.

EXAMPLE

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

clld(1), clload(3), stdcl(3)

libstdcl-1.2

2011-6-13

CLOPEN(3)

CLMALLOC(3)

Standard Compute Layer (CL) Manual

CLMALLOC(3)

NAME

clmalloc, clfree, clsizeofmem - Allocate and free dynamic memory with CL bindings for use with co-processor devices

SYNOPSIS

```
#include <stdcl.h>
```

```
void* clmalloc( CLCONTEXT* cp, size_t size, int flags);
```

```
void* clmrealloc( CLCONTEXT* cp, void* ptr, size_t size, int flags);
```

```
void clfree( void* ptr);
```

```
void clmctl( void* ptr);
```

```
int clmctl( void* ptr, int op, ... );
```

```
int clmctl_va( void* ptr, int op, va_list ap);
```

```
size_t clsizeofmem(void* ptr);
```

Link with -lstdcl.

DESCRIPTION

clmalloc() allocates memory suitable for sharing between compute layer (CL) co-processor devices within a CL context. clmalloc() allocates

size bytes and returns a pointer to the allocated memory. The memory is not cleared. If size is 0, then `clmalloc()` returns a unique pointer value that can later be safely passed to `clfree()`. If called with the flag `CL_MEM_DETACHED` the allocation will no be attached to a CL context.

`clmrealloc()` re-allocates memory suitable for sharing between compute layer (CL) co-processor devices within a CL context. If ptr is 0 or NULL `clmalloc()` is called. If size is 0, then `clmalloc()` returns a unique pointer value that can later be safely passed to `clfree()`.

`clfree()` frees the memory space pointed to by ptr, which must have been returned by a previous call to `clmalloc()`. Otherwise, or if `clfree(ptr)` has already been called before, the behavior is undefined. It is considered an error to call `clfree(ptr)` if ptr is 0 or NULL.

`clmctl()` and `clmctl_va()` are used to control the state of the allocated memory. The only distinction is that `clmctl_va()` accepts a `va_list` directly.

`clsizeofmem()` returns the size of the allocated memory associated with ptr. If ptr does not reference memory allocated by a call to `clmalloc()`, and for which `clfree()` has not been called, the behavior is undefined.

RETURN VALUE

If successful `clmalloc(3)` returns a pointer to the allocated memory that is suitably aligned and suitable for sharing with CL co-processor devices. On error, returns NULL.

`clfree()` returns no value.

`clsizeofmem()` returns the size in bytes of the memory pointed to by ptr.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

`clmattach(3)`, `clmdetach(3)`, `clmsync(3)`, `stdcl(3)`, `malloc(3)`

NAME

clmattach, clmdetach - Attach and detach memory from a CL context

SYNOPSIS

```
#include <stdcl.h>
```

```
int clmattach( CLCONTEXT* cp, void* ptr );
```

```
int clmdetach( void* ptr );
```

Link with -lstdcl.

DESCRIPTION

clmattach() is used to attach memory to a compute layer (CL) context. The memory pointed to by ptr must be allocated with clmalloc() and suitable for sharing between the host and CL co-processor devices. In order to change the attachment of memory from one CL context to another, the memory must first be unattached using a call to clmdetach(). It is an error to pass clmattach() memory that is already attached to a CL context.

clmdetach() is used to detach memory from a CL context. The memory pointed to by ptr must be allocated with clmalloc() and suitable for sharing between the host and CL co-processor devices.

If ptr does not point to memory allocated by clmalloc() the behavior of clmattach() and clmdetach() is undefined.

RETURN VALUE

Both clmattach() and clmdetach() return 0 on success. On error, -1 is returned and errno is set appropriately.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

clmalloc(3), clfree(3), clmsync(3), malloc(3), stdcl(3)

NAME

clmsync - Synchronize memory between host and co-processor device

SYNOPSIS

```
#include <stdcl.h>
```

```
cl_event clmsync( CONTEXT* cp, unsigned int devnum, void* ptr, int flags);
```

Link with -lstdcl.

DESCRIPTION

clmsync() is used to synchronize memory between the host and a compute layer (CL) co-processor device. The memory pointed to by ptr must have been created using a call to clmalloc() and associated with a CL context.

The behavior of clmsync() is controlled by the flags argument which must be set with either CL_MEM_HOST or CL_MEM_DEVICE. These flags are mutually exclusive and it is an error to set both or none. The following flags may be used:

CL_MEM_HOST

clmsync() will sync the memory on the host.

CL_MEM_DEVICE

clmsync() will sync the memory on the device.

CL_EVENT_WAIT

clmsync() will block until the operation has completed.

CL_EVENT_NOWAIT

clmsync() will return immediately. The programmer must ensure that the operation has completed using clwait() or clwaitev().

CL_EVENT_NORELEASE

Used with CL_EVENT_WAIT to prevent clfork() from releasing the CL event generated by the operation. If this flag is used the programmer is responsible for releasing the returned event using clReleaseEvent(). This flag has no effect when CL_EVENT_NOWAIT is used. CL_EVENT_RELEASE This flag is deprecated and should no longer be used.

RETURN VALUE

On error clmsync() will return (cl_event)(-1) and errno is set appropriately.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

`clwait(3)`, `clwaitev(3)`, `clmalloc(3)`, `clfree(3)`, `stdcl(3)`

libstdcl-1.2

2011-6-13

CLMSYNC(3)

CLNDRANGE_INIT(3)

Standard Compute Layer (CL) Manual

CLNDRANGE_INIT(3)

NAME

`clndrange_init1d`, `clndrange_init2d`, `clndrange_init3d` - Initialize the index-space (NDRange) for the execution of a CL kernel

SYNOPSIS

```
#include <stdcl.h>
```

```
clndrange_t clndrange_init1d( gtoff0,gt0,lt0);
```

```
clndrange_t clndrange_init2d( gtoff0,gt0,lt0, gtoff1,gt1,lt1);
```

```
clndrange_t clndrange_init3d( gtoff0,gt0,lt0, gtoff1,gt1,lt1,
gtoff2,gt2,lt2);
```

DESCRIPTION

`clndrange_init()` family of macros are used to initialize an object of type `clndrange_t` that defines the index-space for the execution of a CL kernel. The values of `gtoffn`, `gtn`, `ltn` define the global index offset, global index range and local index range, respectively, for dimension `n`. The index-space defines the work-group and work-item partitioning for the kernel execution.

EXAMPLES

The initialization of a 1-D index-space of 16 work-items with work-group size of 2 and no global offset:

```
clndrange_t ndr = clndrange_init1d( 0,16,2 );
```

The initialization of a 2-D index-space of 64 by 128 work-items with work-group size of 2 by 4 with a global work-item offset of 32,64:

```
clndrange_t ndr = clndrange_init1d( 32,64,2, 64,128,4 );
```

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

clndrange_set(3), clfork(3), stdcl(3)

libstdcl-1.2

2011-6-13

CLNDRANGE_INIT(3)

CLNDRANGE_SET(3)

Standard Compute Layer (CL) Manual

CLNDRANGE_SET(3)

NAME

clndrange_set1d, clndrange_set2d, clndrange_set3d - Set the index-space (NDRange) for the execution of a CL kernel

SYNOPSIS

```
#include <stdcl.h>
```

```
clndrange_set1d( clndrange_t ndr, gtoff0,gt0,lt0);
```

```
clndrange_set2d( clndrange_t ndr, gtoff0,gt0,lt0, gtoff1,gt1,lt1);
```

```
clndrange_set3d( clndrange_t ndr, gtoff0,gt0,lt0, gtoff1,gt1,lt1,
gtoff2,gt2,lt2);
```

DESCRIPTION

clndrange_set() family of macros are used to set an object of type clndrange_t to define the index-space for the execution of a CL kernel. The values of gtoffn, gtn, ltn define the global index offset, global index range and local index range, respectively, for dimension n. The index-space defines the work-group and work-item partitioning for the kernel execution.

EXAMPLES

Setting a 1-D index-space of 16 work-items with work-group size of 2 and no global offset:

```
clndrange_t ndr;
...
clndrange_set1d( ndr, 0,16,2 );
```

Setting a 2-D index-space of 64 by 128 work-items with work-group size of 2 by 4 with a global work-item offset of 32,64:

```
clndrange_t ndr;
...
clndrange_set2d( ndr, 32,64,2, 64,128,4 );
```

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

clndrange_init(3), clfork(3), stdcl(3)

libstdcl-1.2

2011-6-13

CLNDRANGE_SET(3)

CLARG_SET(3)

Standard Compute Layer (CL) Manual

CLARG_SET(3)

NAME

clarg_set, clarg_set_global, clarg_set_local - Set CL kernel arguments

SYNOPSIS

```
#include <stdcl.h>
```

```
void clarg_set( CLCONTEXT* cp, cl_kernel krn, unsigned int argnum, Tn  
arg);
```

```
void clarg_set_global( CLCONTEXT* cp, cl_kernel krn, unsigned int  
argnum, void* ptr);
```

```
void clarg_set_local( CLCONTEXT* cp, cl_kernel krn, unsigned int  
argnum, size_t sizeb);
```

DESCRIPTION

clarg_set(), clarg_set_global() and clarg_set_local() are used to set the argnum argument of the CL kernel krn prior to kernel execution.

clarg_set() is used for setting arguments of intrinsic type such as cl_int, cl_float or cl_float4, etc. For clarg_set() Tn can be any valid scalar or vector type.

clarg_set_global is used for setting arguments of pointers to global memory where ptr points to memory that was allocated using a call to clmalloc() and attached to the CL context of the target kernel.

clarg_set_local() is used for setting arguments of pointers to local memory where sizeb indicates the size in bytes of the local memory that is to be allocated.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

clfork(3), clsym(3), clmalloc(3), stdcl(3)

libstdcl-1.2

2011-6-13

CLARG_SET(3)

CLFORK(3)

Standard Compute Layer (CL) Manual

CLFORK(3)

NAME

clfork - Execute a CL kernel

SYNOPSIS

```
#include <stdcl.h>
```

```
cl_event clfork( CLCONTEXT* cp, unsigned int devnum, cl_kernel krn,
clndrange_t* ndr, int flags);
```

Link with -lstdcl.

DESCRIPTION

clfork() is used to execute a CL kernel on a specified compute layer (CL) co-processor device. The arguments for the kernel must be set prior to the call to clfork() using the clarg_set*() functions. The kernel is executed over an index-space of work-items defined by ndr.

The behavior of clfork() can be controlled using the following flags:

CL_EVENT_WAIT

clfork() will block until the operation has completed.

CL_EVENT_NOWAIT

clfork() will return immediately. The programmer must ensure that the operation has completed using clwait() or clwaitev().

CL_EVENT_NORELEASE

Used with CL_EVENT_WAIT to prevent clfork() from releasing the CL event generated by the operation. If this flag is used the programmer is responsible for releasing the returned event using clReleaseEvent(). This flag has no effect when CL_EVENT_NOWAIT is used. CL_EVENT_RELEASE This flag is deprecated and should no

longer be used.

RETURN VALUE

On error `clfork()` will return `(cl_event)(-1)` and `errno` is set appropriately.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

`clarg_set(3)`, `clndrange_init(3)`, `clndrange_set(3)`, `clwait(3)`,
`clwaitev(3)`, `stdcl(3)`

libstdcl-1.2

2011-6-13

CLFORK(3)

CLFLUSH(3)

Standard Compute Layer (CL) Manual

CLFLUSH(3)

NAME

`clflush` - Flush the CL command queue

SYNOPSIS

```
#include <stdcl.h>
```

```
int clflush( CLCONTEXT* cp, cl_uint devnum, int flags);
```

Link with `-lstdcl`.

DESCRIPTION

`clflush()` is used to flush the OpenCL command queue for device number `devnum` within a CL context. For certain OpenCL implementations this is necessary to initiate operations to be executed asynchronously.

The `flags` argument is reserved for future use and presently ignored.

RETURN VALUE

On error `clflush()` will return `(cl_event)(-1)` and `errno` is set appropriately.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to <support@browndeertechnology.com>

COPYRIGHT

Copyright (C) 2010-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

clfork(3), clmsync(3), clwait(3), stdcl(3)

libstdcl-1.2

2011-6-13

CLFLUSH(3)

CLWAIT(3)

Standard Compute Layer (CL) Manual

CLWAIT(3)

NAME

clwait - Block on one or more CL events

SYNOPSIS

```
#include <stdcl.h>
```

```
cl_event clwait( CLCONTEXT* cp, cl_uint devnum, int flags);
```

Link with -lstdcl.

DESCRIPTION

clwait() is used to block on the completion of one or more outstanding events for device number devnum within a CL context. The type of events are specified by selecting one or more event lists as described below.

One or more event lists may be selected using a combination of the following flags:

CL_KERNEL_EVENT

Block on events in the ordered kernel event list.

CL_MEM_EVENT

Block on events in the ordered memory event list.

Note that if both kernel and memory event lists are specified, the kernel event list has first priority. Specifically, clwait() will first block on all outstanding kernel events and subsequently block on all outstanding memory events.

The behavior of clwait() can be controlled using the following flags:

CL_EVENT_NORELEASE

Used with CL_EVENT_WAIT to prevent clfork() from releasing the CL event generated by the operation. If this flag is used the programmer is responsible for releasing the returned event using

`clReleaseEvent()`. This flag has no effect when `CL_EVENT_NOWAIT` is used. `CL_EVENT_RELEASE` This flag is deprecated and should no longer be used.

RETURN VALUE

On error `clwait()` will return `(cl_event)(-1)` and `errno` is set appropriately.

AUTHOR

Written by David Richie.

REPORTING BUGS

Report bugs to [<support@browndeertechnology.com>](mailto:support@browndeertechnology.com)

COPYRIGHT

Copyright (C) 2009-2011 Brown Deer Technology, LLC. Licensed under the GNU Lesser General Public License version 3 (LGPLv3). There is NO WARRANTY to the extent permitted by law.

SEE ALSO

`clfork(3)`, `clmsync(3)`, `clwaitev(3)`, `stdcl(3)`